

Bayesian inference using MCMC in OpenTURNS

E. Songo ¹ J. Muré ¹ M. Keller ¹

¹EDF R&D. 6, quai Watier, 78401, Chatou Cedex - France, joseph.mure@edf.fr

May 20th 2024, BEPU 2024, Lucca (Italy)



OpenTURNS: www.openturns.org

► Data analysis

- Distribution fitting
- Statistical tests
- Estimate dependency and copulas
- Estimate stochastic processes

► Probabilistic modeling

- Dependence modeling
- Univariate distributions
- Multivariate distributions
- Copulas
- Processes
- Covariance kernels

► Surrogate models

- Linear regression
- Polynomial chaos expansion
- Gaussian process regression
- Spectral methods
- Low rank tensors
- Fields metamodel

► Reliability, sensitivity





- Sampling methods
- Approximation methods
- Sensitivity analysis
- Design of experiments

► Calibration

- Least squares calibration
- Gaussian calibration
- Bayesian calibration


► Numerical methods

- Optimization
- Integration
- Least squares
- Meshing
- Coupling with external codes



Coupling OpenTURNS with computer codes

OpenTURNS provides a text file exchange based interface in order to perform analyses on complex computer codes



- ▶ Replaces the need for input/output text parsers
- ▶ Wraps a simulation code under the form of a standard python function
- ▶ Allows to interface OpenTURNS with a cluster
- ▶ `otwrapy`: interfacing tool to allow easy parallelization

Contents

About OpenTURNS

The Metropolis-Hastings algorithm

Random walk Metropolis-Hastings

The Gibbs algorithm

Independent Metropolis-Hastings

User-defined Metropolis-Hastings

Application

Conclusion

OpenTURNS: Metropolis-Hastings

We want to sample from the distribution π of a random variables X . Here is one step of the algorithm, starting from the point x :

1. Simulate a candidate $x' \sim q(\cdot|x)$ for some conditional distribution q .
2. Compute $\alpha(x'|x, y, z) = \min \left\{ \frac{\pi(x')q(x|x')}{\pi(x)q(x'|x)}, 1 \right\}$.
3. Simulate $u \sim \mathcal{U}(0, 1)$. If $u \leq \alpha(x'|x)$, then the next state is x' , otherwise it is x .


Throughout the presentation, our code is prefaced by:

```
import openturns as ot
import math as m
import numpy as np
```

Random walk Metropolis Hastings

When $q(\cdot|x) = x + \mu$, where μ is a distribution that does not depend on x , the algorithm is called “Random walk Metropolis-Hastings” and μ is called the “proposal distribution”.

Sample from a nonstandard distribution¹




$$\begin{aligned} \pi(x) \\ \propto \frac{1}{2} (2 + \sin(x)^2) \\ \exp \left[- \left(2 + \cos(3x)^3 + \sin(2x)^3 \right) x \right] \\ \mathbf{1}_{[0,2\pi]}(x) \end{aligned}$$

```
logdensity = ot.SymbolicFunction('x', 'log(2+sin(x)^2) - (2+cos(3*x)^3+sin(2*x)^3) * x')
support = ot.Interval([0.0], [2.0 * m.pi])
proposal = ot.Normal(0.0, 2.0) # mu
initialState = [3.0]
sampler = ot.RandomWalkMetropolisHastings(logdensity, support, initialState, proposal)
x = sampler.getSample(10000)
```

¹Marin, J.M. and Robert, C.P. (2007). Bayesian Core: A Practical Approach to Computational Bayesian Statistics. Springer-Verlag, New York

2D Random walk Metropolis Hastings

Sample from a 2D nonstandard distribution




$$\begin{aligned} \pi(x) &\propto \left(\exp \left[-\frac{1}{4}(x-3)^2 + y^2 \right] \right. \\ &\quad \left. + \exp \left[-(x-5)^2 - 5 \left(y - \frac{1}{5} \right)^2 \right] \right) \\ &\quad \mathbf{1}_{[0,2\pi]}(x) \mathbf{1}_{[0,1]}(y) \end{aligned}$$

```
logdensity = ot.SymbolicFunction(
    ["x", "y"], ["log(exp(-0.25 * (x-3)^2 + y^2) + exp(-(x-5)^2 - 5 * (y-0.2)^2))"]
)
support = ot.Interval([0.0, 0.0], [2.0 * m.pi, 1.0])
proposal = ot.Normal([0.0] * 2, [1.0, 0.2])
initialState = [3.0, 0.8]
sampler = ot.RandomWalkMetropolisHastings(logdensity, support, initialState, proposal)
x = sampler.getSample(50000)
```

2D Random walk Metropolis Hastings in a Bayesian setting

Posterior distribution of the parameters of a Weibull model



$$\beta \sim \Gamma(k = 2, \lambda = 2 \cdot 10^{-4})$$

$$\alpha \sim \mathcal{U}(0.5, 3.8)$$

$$T|\beta, \alpha \sim \mathcal{W}(\beta, \alpha, 0)$$


$$F_{\mathcal{W}}(t) = 1 - \exp \left[- \left(\frac{t - 0}{\beta} \right)^{\alpha} \right]$$

```
alpha_min, alpha_max, a_beta, b_beta = 0.5, 3.8, 2.0, 2.0e-4
priorMarginals = [ot.Gamma(a_beta, b_beta), ot.Uniform(alpha_min, alpha_max)]
prior = ot.ComposedDistribution(priorMarginals)
proposal = ot.Normal([0.0]*2, [0.1*m.sqrt(a_beta/b_beta**2), 0.1*(alpha_max-alpha_min)])
initialState = [a_beta / b_beta, 0.5 * (alpha_max - alpha_min)]
sampler = ot.RandomWalkMetropolisHastings(prior, initialState, proposal)

conditional = ot.WeibullMin()
Tobs = [[4380], [1791], [1611], [1291]]

# WeibullMin expects beta, alpha, and localization, but the prior is only on beta, alpha
linkFunction = ot.SymbolicFunction(["beta", "alpha"], ["beta", "alpha", "0"])
sampler.setLikelihood(conditional, Tobs, linkFunction)
sample = sampler.getSample(100000)
```


A flood model



$$\forall 1 \leq i \leq 8, H^{(i)} \sim$$

$$\mathcal{N}\left(G(Q^{(i)}, K_s, Z_v, Z_m), \frac{1}{2}\right)$$

$$K_s \sim \mathcal{N}(20, 5)$$

$$Z_v \sim \mathcal{N}(49, 1)$$

$$Z_m \sim \mathcal{N}(51, 1)$$

```
Qobs = [[2097], [1448], [1516], [2173], [387], [3016], [651], [541]]
Hobs = [[3.4], [2.5], [2.7], [3.5], [1.0], [4.2], [1.6], [1.6]]
```

```
def flooding(X):
    L = 5.0e3
    B = 300.0
    Q, K_s, Z_v, Z_m = X
    alpha = (Z_m - Z_v) / L
    if alpha < 0.0 or K_s <= 0.0:
        H = np.inf
    else:
        H = (Q / (K_s * B * np.sqrt(alpha))) ** (3.0 /
            5.0)
    return [H, 0.5]
```

```
functionG = ot.PythonFunction(4, 2, flooding)
```


```
# Q (input #0) is not calibrated
linkFunction = ot.ParametricFunction(functionG, [0], [100])
```

```
conditional = ot.Normal()
```

```
parameterPriorMean = [20.0, 49.0, 51.0]
parameterPriorSigma = [5.0, 1.0, 1.0]
prior = ot.Normal(parameterPriorMean, parameterPriorSigma)
```

```
initialState = parameterPriorMean
```

Single component Random Walk Metropolis-Hastings




```


mh_coll = [
    ot.RandomWalkMetropolisHastings(prior, initialState, ot.Uniform(-5.0, 5.0), [0]),
    ot.RandomWalkMetropolisHastings(prior, initialState, ot.Uniform(-1.0, 1.0), [1]),
    ot.RandomWalkMetropolisHastings(prior, initialState, ot.Uniform(-1.0, 1.0), [2]),
]

for mh in mh_coll:
    mh.setLikelihood(conditional, Hobs, linkFunction, Qobs)


sampler = ot.Gibbs(mh_coll) # NB: the order can be made random: cf. setUpdatingMethod
sample = sampler.getSample(10000)
  
```



Muré (EDF)



OpenTURNS




Blocks of components can be considered




```


mh_coll = [
    ot.RandomWalkMetropolisHastings(prior, initialState, ot.Uniform(-5.0, 5.0), [0]),
    ot.RandomWalkMetropolisHastings(prior,
                                    initialState,
                                    ot.ComposedDistribution([ot.Uniform(-1.0,1.0)]*2),
                                    [1, 2])
]
for mh in mh_coll:
    mh.setLikelihood(conditional, Hobs, linkFunction, Qobs)
sampler = ot.Gibbs(mh_coll) # NB: the order can be made random: cf. setUpdatingMethod
sample = sampler.getSample(10000)
  
```



Muré (EDF)




OpenTURNS



Independent Metropolis-Hastings: $q(\cdot|x) = \mu$

Instrumental PDF




```
logdensity = ot.SymbolicFunction('x','...') # replace ...
support = ot.Interval([0.0], [2.0 * m.pi])
initialState = [3.0] # unimportant for independent MH


exp = ot.Exponential(1.0)
unif = ot.Normal(5.3, 0.4)
instrumental = ot.Mixture([exp, unif], [0.9, 0.1])

independentMH = ot.IndependentMetropolisHastings(
    logdensity, support, initialState, instrumental
)
x = independentMH.getSample(10000)
```

Random walk Metropolis-Hastings




Independent Metropolis-Hastings



User-defined Metropolis-Hastings: $q(\cdot|x) = \mu(x)$

Metropolis adjusted Langevin algorithm² implementation



With $h > 0$ a fixed parameter:

$$q(\cdot|x) = \mathcal{N}\left(x + \frac{h}{2} \frac{d}{dx} [\log(\pi(x))], \sqrt{h}\right)$$

```
from openturns.experimental import UserDefinedMetropolisHastings
logdensity = ot.SymbolicFunction('x', 'log(2+sin(x)^2) - (2+cos(3*x)^3+sin(2*x)^3) * x')
support, proposal, initialState = ot.Interval([0.0], [2.0 * m.pi]), ot.Normal(), [2.5]
h = 0.5
std_deviation = m.sqrt(h)

def python_link(x):
    derivative_log_density = logdensity.getGradient().gradient(x)[0, 0]
    mean = x[0] + h / 2 * derivative_log_density
    return [mean, std_deviation]
link = ot.PythonFunction(1, 2, python_link)

mala = UserDefinedMetropolisHastings(logdensity, support, initialState, proposal, link)
z = mala.getSample(10000)
```

²Robert, C. P. *The Metropolis-Hastings algorithm*. arXiv preprint arXiv:1504.01896, 2015

Application: airflow rate in a depressurized room

$$g(\xi, \theta_0, \theta_1) = 0.6 \times 3600 \times \theta_0 \left(\frac{2}{1.8} \xi \right)^{\theta_1}$$

For $1 \leq i \leq 233$:

$$\begin{cases} X_i &= \xi_i + Z_i \\ Y_i &= g(\xi_i, \theta_0, \theta_1) + E_i \end{cases}$$

► Input: X_i (pressure difference – bars)

► Output: Y_i (airflow rate – m^3/h)

► Parameter: θ_0 (area – m^2)

► Parameter: θ_1 (exponent)

► $Z_i \sim U(-0.05, 0.05)$

► $\theta_0 \sim U(0, 2)$










► $E_i \sim N(0, \sigma_E^2), \sigma_E^2 \sim 1/\sigma_E^2$

► $\theta_1 \sim U(0, 2)$

Strategy: σ_E^2 averaged out analytically, the rest sampled using Gibbs with:

- Random walk Metropolis-Hastings on θ_0 , step is tuned during burn-in.
- Random walk Metropolis-Hastings on θ_1 , step is tuned during burn-in.
- Independent MH on each Z_i with the prior as proposal.

Convergence diagnostics for θ_0 , θ_1 , ξ_1 with 3 chains



(a) ACF plot

(b) Trace plot

(c) Convergence plot of ergodic means

Posterior distribution

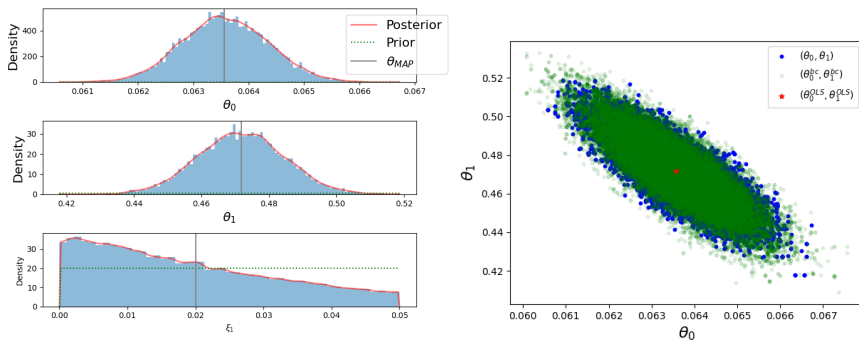


Figure: Left: Prior and posterior distributions of θ_0 , θ_1 and ξ_1 . Right: Scatter plot of the sample of (θ_0, θ_1) from the joint posterior distribution (solid blue dots), from a simplified posterior where all $Z_i = 0$ (transparent green dots), alongside the Ordinary Least Squared estimator (red star)

Conclusion

OpenTURNS provides an MCMC sampling framework through the following classes:

- ▶ MetropolisHastings variants:
 - ▶ RandomWalkMetropolisHastings
 - ▶ IndependentMetropolisHastings
 - ▶ UserDefinedMetropolisHastings
 - ▶ RandomVectorMetropolisHastings (not shown in this presentation)
- ▶ Gibbs

These classes can be freely combined to sample from nonstandard distributions in a “smart” manner.

In a Bayesian setting, this framework allows users to create and implement the MCMC algorithm most suited to a particular posterior distribution.